

Name: Hao Ren
NetID: haor2
Section: ZJ1

ECE 408/CS483 Milestone 3 Report

0. List Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images from your basic forward convolution kernel in milestone 2. This will act as your baseline this milestone. Note: **Do not** use batch size of 10k when you profile in `--queue rai_amd64_exclusive`. We have limited resources, so any tasks longer than 3 minutes will be killed. Your baseline M2 implementation should comfortably finish in 3 minutes with a batch size of 5k (About 1m35 seconds, with nv-nsight).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.246515 ms	0.83401 ms	0m1.542s	0.86
1000	2.24226 ms	15.3607 ms	0m10.669s	0.886
5000	11.5922 ms	40.737 ms	0m0.740s	0.871

1. **Optimization 1: `IMPL_INPUT_UNROLLING` Shared memory matrix multiplication and input matrix unrolling (**3 points**)**

- a. Which optimization did you choose to implement and why did you choose that optimization technique.

The current access pattern and techniques are convolution which isn't optimized and doesn't have special hardware support. To further optimize the performance, we can first convert the whole computing pattern to matrix multiplication, hence a conv2mul kernel is needed to unroll the input matrix.

The expected improvement might not be that much (or degrades the performance), but it leaves room for off-load the workload to tensor core.

- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

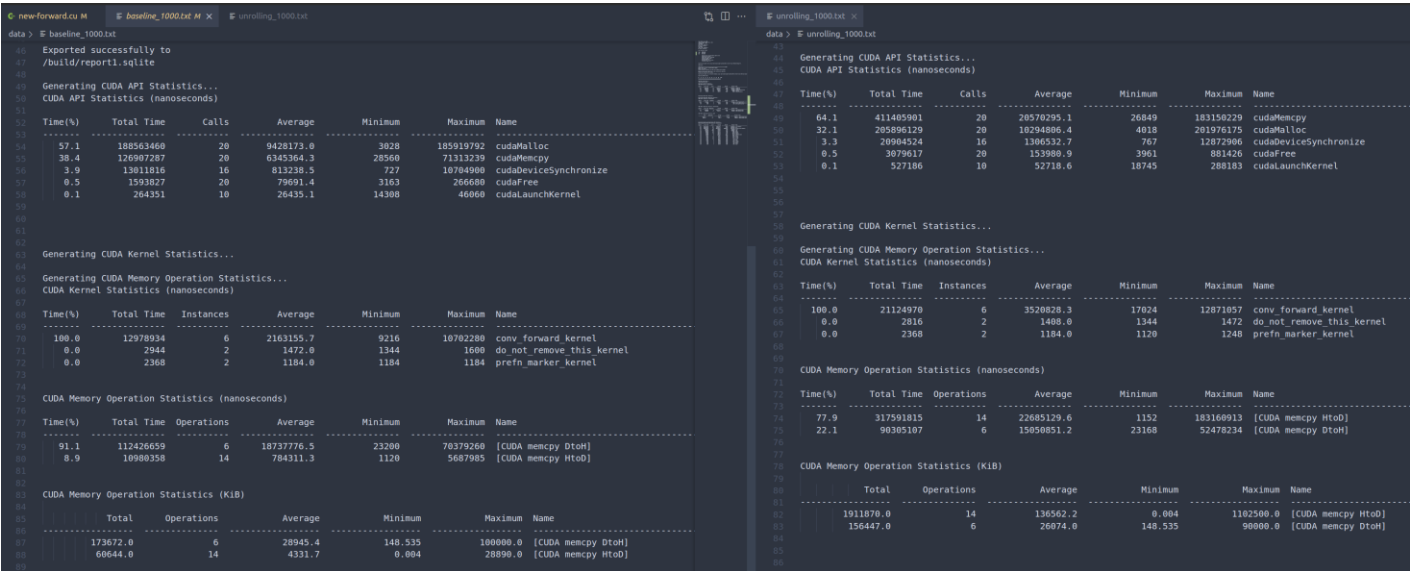
The optimization works by changing the computing method from convolution to matrix multiplication. NO, I don't think it will improve the performance at this point since the unrolling can be expensive at this point and memory coalescing also happens in baseline when set properly. Since this is the first optimization, I will talk about the synergize in next optimization.

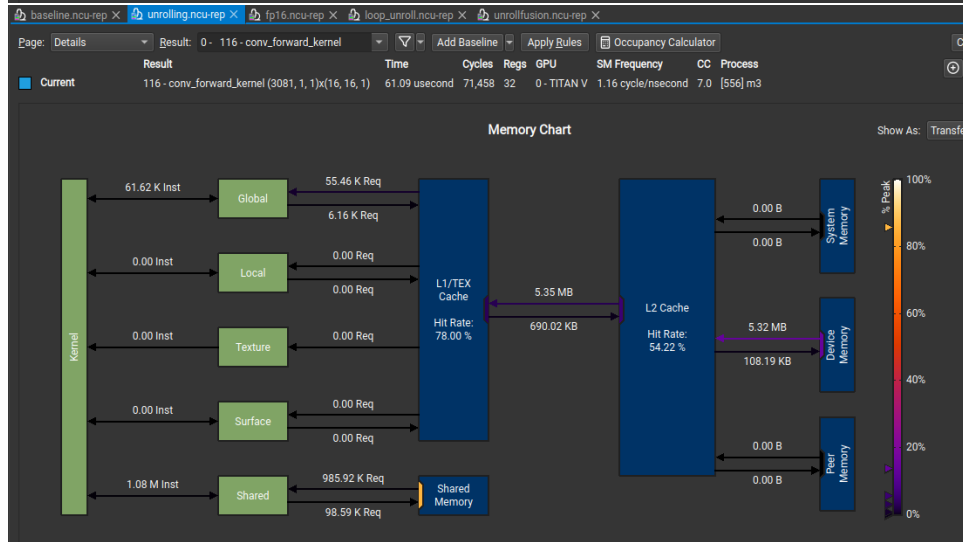
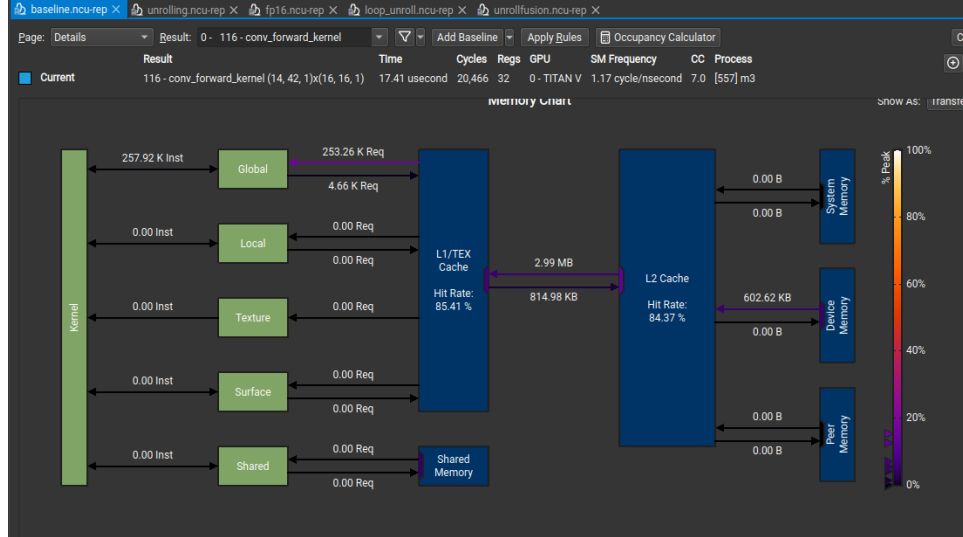
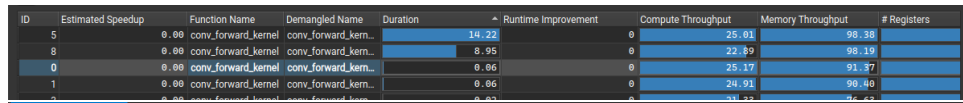
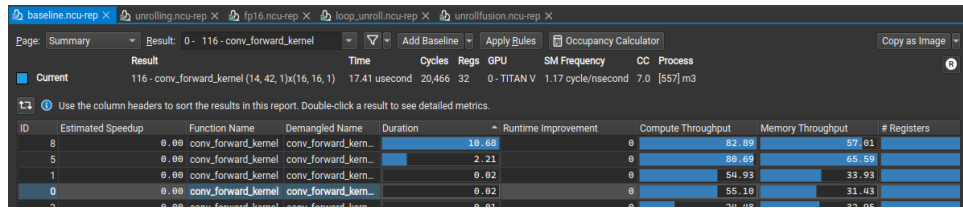
- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	1.45413 ms	0.939222 ms	0m0.297s	0.86
1000	14.3275 ms	9.03383 ms	0m1.064s	0.886
5000	71.0609 ms	44.6977 ms	0m3.832s	0.871

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from nsys and Nsight-Compute to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).

NSYS





No. Because the convolution kernel, as expected, takes longer time due to the fact that we have to load the data into shared memory first, and the copying between shared memory and global memory takes time even longer than pure global memory access due to the fact that we have better memory coalescing in pure global memory access pattern because of my design. The exec time will also be larger because the input unroll makes the memory usage about $K \times K$ times larger than baseline. (As we can see from nsys, cudaMalloc and cudaMemcpy takes significantly longer time, conv_forward_kernel also takes longer time).

*When we compare nsight compute GUI, we can see this optimization uses more memory bandwidth because we have much larger input because of input unroll ($K*K$ larger, around 49 times)*

- e. What references did you use when implementing this technique?

The lecture and MP4. Copilot also helps.

2. Optimization 2: [IMPL_UNROLLING_KERNEL_FUSION](#) Kernel fusion for unrolling and matrix-multiplication (requires previous optimization) (2 points**)**

- a. Which optimization did you choose to implement and why did you choose that optimization technique.

I use kernel fusion for the original CPU unrolling to kernel input unrolling and I optimized the original matrix multiplication.

- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

The optimization works by running input unrolling and finding a better way to do kernel multiplication. I think this will improve the op time as well as the exec time as the work was done in CPU in the previous optimization now it's done in GPU, and I tried the way with better memory coalescing. This optimization is built on previous optimization which gives a baseline of input unroll version implementation. In this optimization, the performance should be significantly improved.

- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.234637 ms	1.08158 ms	0m0.181s	0.86
1000	2.24452 ms	10.7243 ms	0m0.332s	0.886
5000	11.0265 ms	53.3184 ms	0m0.925s	0.871

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).

new-forward M

unrolling_1000.dat X

re_build.yml M

data > E unrolling_1000.dat

Exported successfully to /build/report1.sqlite

Generating CUDA API Statistics... CUDA API Statistics (nanoseconds)

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
44.1	43148981	20	2057925.1	26449	18318829	cudaMemcpy
32.1	28596129	20	1029486.4	4818	201976175	cudaMalloc
3.3	29964524	16	136632.7	767	12872906	cudaDeviceSynchronize
0.5	3076017	20	15388.9	3961	851426	cudaFree
0.1	527186	10	52718.6	18745	288163	cudaLaunchKernel

Generating CUDA Kernel Statistics... Generating CUDA Memory Operation Statistics... CUDA Kernel Statistics (nanoseconds)

Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
106.0	21124970	6	3520828.3	17024	12871057	conv_forward_kernel
0.0	2310	2	1480.0	1344	1512	do_not_remove_this_kernel
0.0	2368	2	1184.0	1120	1248	prefn_marker_kernel

CUDA Memory Operation Statistics (nanoseconds)

Time(%)	Total Time	Operations	Average	Minimum	Maximum	Name
77.9	31798835	14	22665129.6	1152	183188913	[CUDA memory H2D]
22.1	98385187	6	15958851.2	23168	52478234	[CUDA memory D2H]

CUDA Memory Operation Statistics (KiB)

Total	Operations	Average	Minimum	Maximum	Name
1811879.0	14	130562.2	0.004	1182488.0	[CUDA memory H2D]
156447.0	6	26074.0	148.535	98000.0	[CUDA memory D2H]

data > E unrollfusion_1000.dat

Exported successfully to /build/report1.sqlite

Generating CUDA API Statistics... CUDA API Statistics (nanoseconds)

Time(%)	Total Time	Calls	Average	Minimum	Maximum	Name
51.9	321659387	22	14620472.1	880	196829661	cudaDeviceSynchronize
28.9	174669974	32	5608155.4	2878	174859832	cudaMalloc
19.1	112530331	20	5626516.5	39327	50894172	cudaMemcpy
1.0	6294420	32	196786.6	3872	1881766	cudaFree
0.1	390644	10	24732.7	18825	42486	cudaLaunchKernel

Generating CUDA Kernel Statistics... Generating CUDA Memory Operation Statistics... CUDA Kernel Statistics (nanoseconds)

Time(%)	Total Time	Instances	Average	Minimum	Maximum	Name
95.3	386360244	6	51866849.7	78943	196885461	in2col_kernel
4.7	15232112	6	253885.3	7392	11184773	conv_forward_kernel
0.0	2816	2	1408.0	1408	1408	do_not_remove_this_kernel
0.0	2560	2	1280.0	1248	1312	prefn_marker_kernel

CUDA Memory Operation Statistics (nanoseconds)

Time(%)	Total Time	Operations	Average	Minimum	Maximum	Name
96.7	98879559	6	16471593.2	23520	57119668	[CUDA memory D2H]
9.3	18092644	14	729514.6	1152	4811319	[CUDA memory H2D]

CUDA Memory Operation Statistics (KiB)

Total	Operations	Average	Minimum	Maximum	Name
173672.0	6	28945.4	148.535	188088.0	[CUDA memory D2H]
68044.0	14	4331.7	0.004	28290.0	[CUDA memory H2D]

baseline.ncu-rep X

unrolling.ncu-rep X

unrollfusion.ncu-rep X

fp16.ncu-rep X

loop_unroll.ncu-rep X

Page: Summary

Result: 0- 116- conv_forward_kernel

Add Baseline

Apply Rules

Occupancy Calculator

Copy as Image

Current

116- conv_forward_kernel (3081, 1, 1)x(16, 16, 1)

61.09 usecond

71,458

32

0- TITAN V

1.16 cycle/usecond

7.0

[556] m3

Use the column headers to sort the results in this report. Double-click a result to see detailed metrics.

ID	Estimated Speedup	Function Name	Demangled Name	Duration	Runtime Improvement	Compute Throughput	Memory Throughput	# Registers
5	0.00	conv_forward_kernel	conv_forward_kern...	14.22		25.01	98.38	
8	0.00	conv_forward_kernel	conv_forward_kern...	8.95		22.89	98.19	
0	0.00	conv_forward_kernel	conv_forward_kern...	25.17		25.17	91.87	
1	0.00	conv_forward_kernel	conv_forward_kern...	0.06		24.91	90.40	
3	0.00	conv_forward_kernel	conv_forward_kern...	0.02		21.33	76.63	
2	0.00	conv_forward_kernel	conv_forward_kern...	0.02		29.43	73.32	
9	0.00	do_not_remove_this...	do_not_remove_thi...	0.00		0.00	0.22	
4	0.00	prefn_marker_kernel	prefn_marker_kern...	0.00		0.00	0.22	
7	0.00	prefn_marker_kernel	prefn_marker_kern...	0.00		0.00	0.22	
6	0.00	do_not_remove_this...	do_not_remove_thi...	0.00		0.00	0.22	

baseline.ncu-rep X

unrolling.ncu-rep X

unrollfusion.ncu-rep X

fp16.ncu-rep X

loop_unroll.ncu-rep X

Page: Summary

Result: 0- 120- in2col_kernel

Add Baseline

Apply Rules

Occupancy Calculator

Copy as Image

Current

120- in2col_kernel (3081, 6, 1)x(16, 16, 1)

301.12 usecond

358,429

30

0- TITAN V

1.19 cycle/usecond

7.0

[556] m3

Use the column headers to sort the results in this report. Double-click a result to see detailed metrics.

ID	Estimated Speedup	Function Name	Demangled Name	Duration	Runtime Improvement	Compute Throughput	Memory Throughput	# Registers
14	0.00	conv_forward_kernel	conv_forward_kern...	11.28		21.31	98.06	
10	0.00	conv_forward_kernel	conv_forward_kern...	3.98		30.16	88.85	
1	0.00	conv_forward_kernel	conv_forward_kern...	0.02		37.07	60.04	
3	0.00	conv_forward_kernel	conv_forward_kern...	0.02		36.46	59.05	
7	0.00	conv_forward_kernel	conv_forward_kern...	0.01		23.68	37.06	
5	0.00	conv_forward_kernel	conv_forward_kern...	0.01		21.68	34.56	
15	0.00	do_not_remove_this...	do_not_remove_thi...	0.00		0.00	0.23	
11	0.00	do_not_remove_this...	do_not_remove_thi...	0.00		0.00	0.22	

baseline.ncu-rep X

unrolling.ncu-rep X

unrollfusion.ncu-rep X

fp16.ncu-rep X

loop_unroll.ncu-rep X

Page: Details

Result: 0- 116- conv_forward_kernel

Add Baseline

Apply Rules

Occupancy Calculator

Copy as Image

Current

116- conv_forward_kernel (3081, 1, 1)x(16, 16, 1)

61.09 usecond

71,458

32

0- TITAN V

1.16 cycle/usecond

7.0

[556] m3

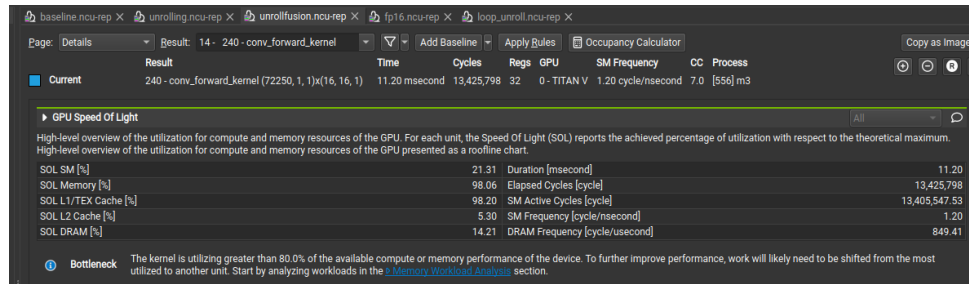
GPU Speed Of Light

High-level overview of the utilization for compute and memory resources of the GPU. For each unit, the Speed Of Light (SOL) reports the achieved percentage of utilization with respect to the theoretical maximum. High-level overview of the utilization for compute and memory resources of the GPU presented as a barline chart.

SOL SM [%]	25.17	Duration [usecond]	61.09
SOL Memory [%]	91.37	Elapsed Cycles [cycle]	71,458
SOL L1/TEX Cache [%]	95.13	SM Active Cycles [cycle]	68,287.96
SOL L2 Cache [%]	5.86	SM Frequency [cycle/usecond]	1.16
SOL DRAM [%]	13.99	DRAM Frequency [cycle/usecond]	827.31

Bottleneck

The kernel is utilizing greater than 80.0% of the available compute or memory performance of the device. To further improve performance, work will likely need to be shifted from the most utilized to another unit. Start by analyzing workloads in the [Memory Workload Analysis](#) section.



Yes. It significantly improves both exec time and Op time. For exec time it's easy because we move CPU exec to GPU kernel exec. For Op time it's actually quite dependent on batch size. When the batch size isn't big enough, using shared memory would add extra overhead since the global memory isn't fully exploited or the over-utilized global memory bandwidth isn't great enough to degrade the performance of kernel fusion in this optimization, from the statistics, it could be seen that although the memory usage (global) in this optimization is greater than the previous optimization, the performance is still better.

- e. What references did you use when implementing this technique?

Asked chatgpt and copilot for some help.

3. Optimization 3: IMPL_LOOP_UNROLL Tuning with restrict and loop unrolling (considered as one optimization only if you do both) (**3 points**)

(Delete this section blank if you did not implement this many optimizations.)

- a. Which optimization did you choose to implement and why did you choose that optimization technique.

I added `__restrict__` keyword to pointer and I also enumerate several cases of K (K=1,2,3,4,7) for unrolling purpose. The restrict pointer can help compiler improve the performance, and the unrolling can significantly improve the baseline approach by improving memory coalescing ability and helping compiler and reduce the use of extra registers which would cause many overheads.

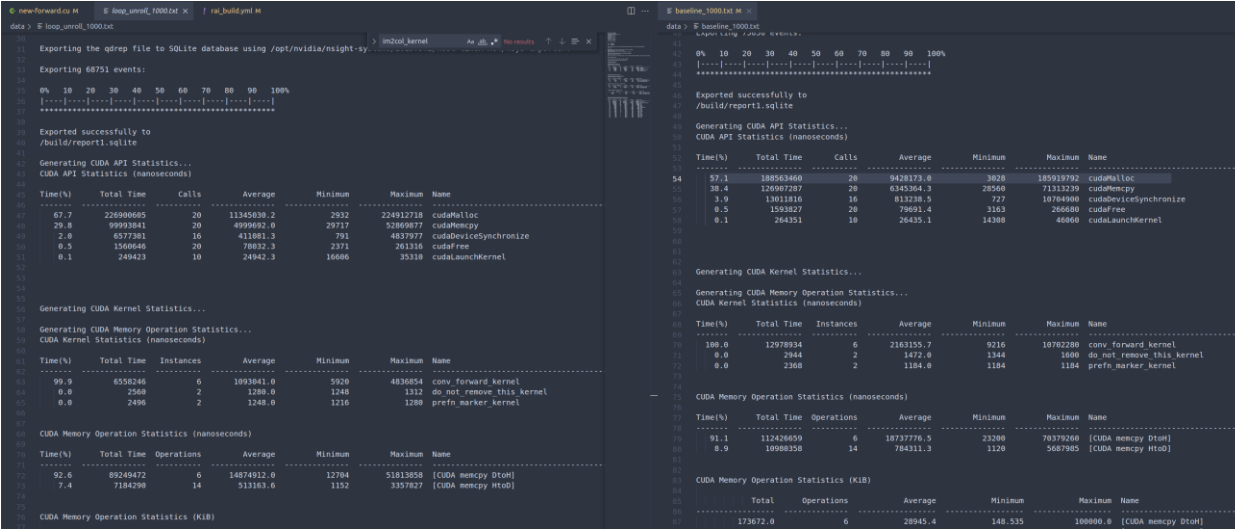
- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

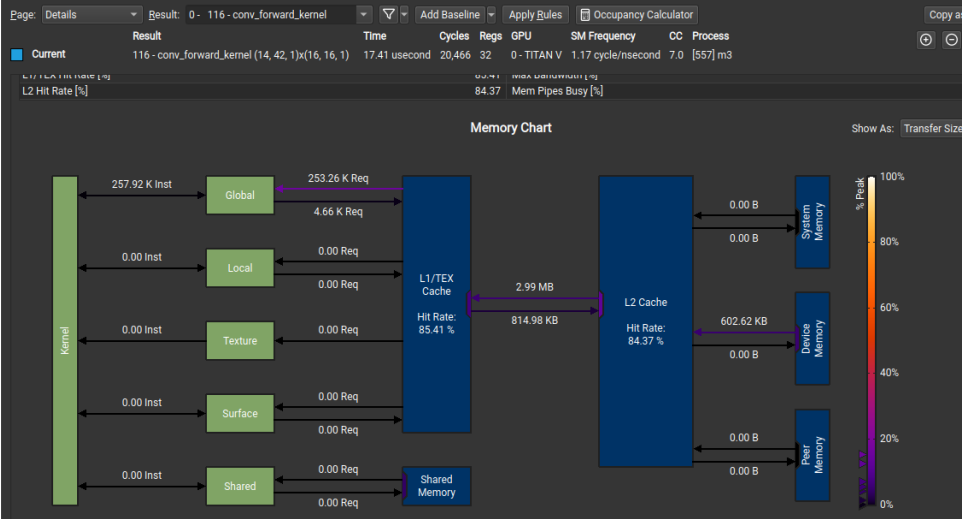
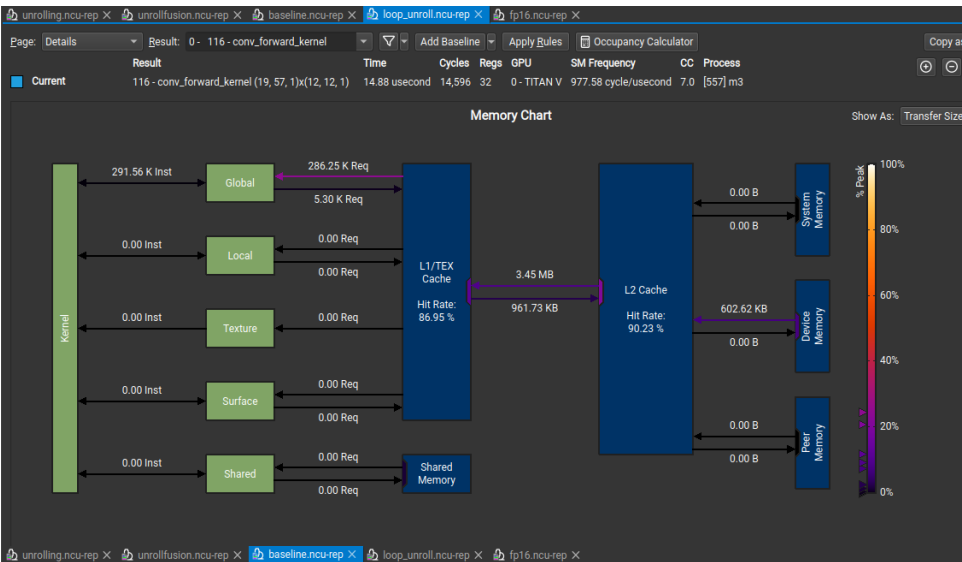
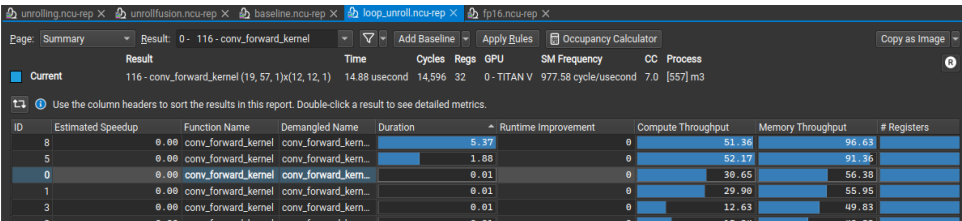
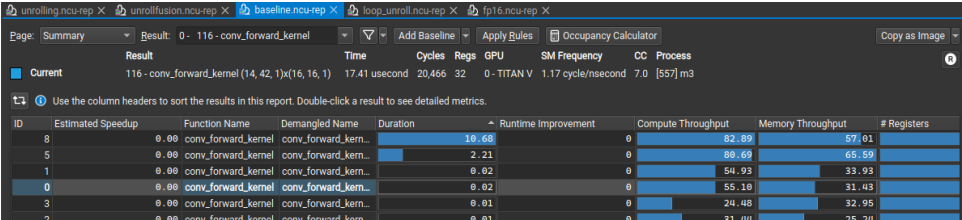
The restrict pointer can help compiler improve the performance by informing the compiler that this memory region is exclusive to this pointer and will not be referenced by other pointer in this kernel, allowing it to make optimization if available, and the unrolling can significantly improve the baseline approach by improving memory coalescing ability and helping compiler and reduce the use of extra registers which would cause many overheads. This would definitely improve the performance for forward convolution for previously stated reasons. This optimization is independent from previous 2 optimizations and should be compared with baseline implementation.

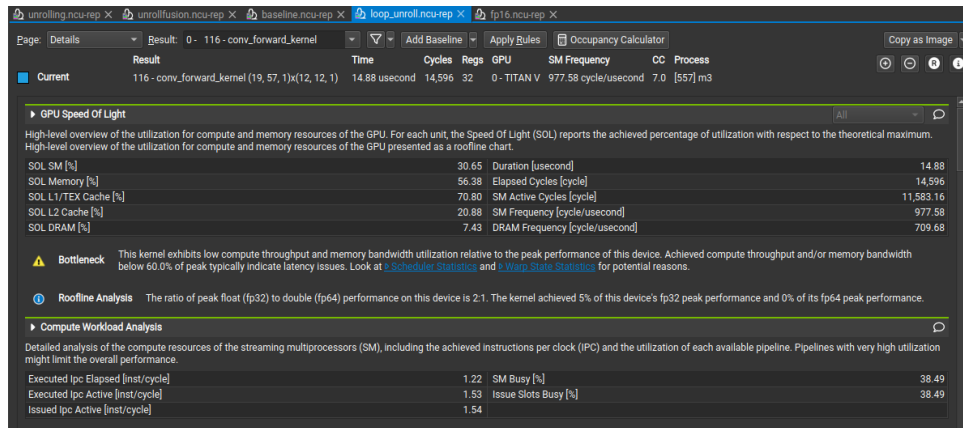
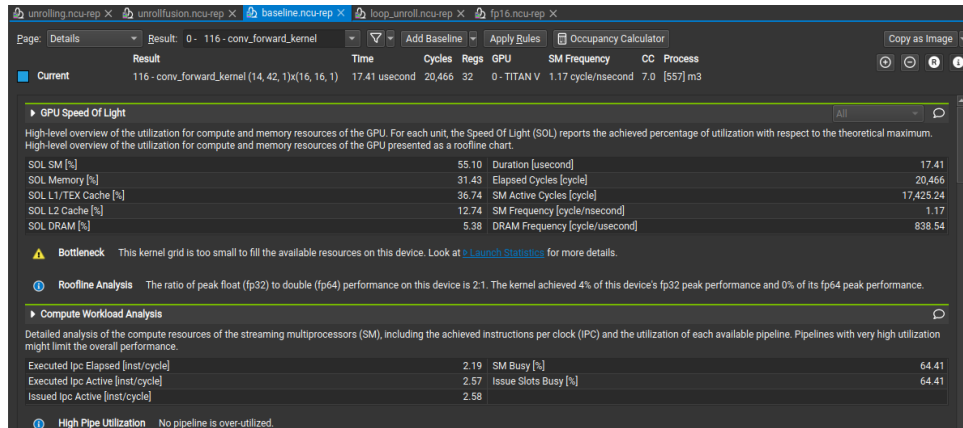
- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.197214 ms	0.540495 ms	0m0.173s	0.86
1000	1.89108 ms	5.41518 ms	0m0.344s	0.886
5000	10.1522 ms	26.9867 ms	0m1.024s	0.871

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from nsys and Nsight-Compute to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off).







The implementation is very successful in terms of performance (Op time) because the memory is better coalesced which can be shown that SOL memory cycles is less in current optimization. However, the total data transferred from memory should (trivially) be unchanged as shown.

- e. What references did you use when implementing this technique?
Chatgpt.

4. Optimization 4: **IMPL_FP16** FP16 arithmetic. (note this can modify model accuracy slightly) (**4 points**) *(Delete this section blank if you did not implement this many optimizations.)*

- a. Which optimization did you choose to implement and why did you choose that optimization technique.

FP16 arithmetic with restrict and loop unrolling because this set of optimizations should be most promising since all the operations can operate on 16bit width operands which will significantly reduce overhead. Although it's note-worthy that the conversion of input and mask and output can introduce overhead.

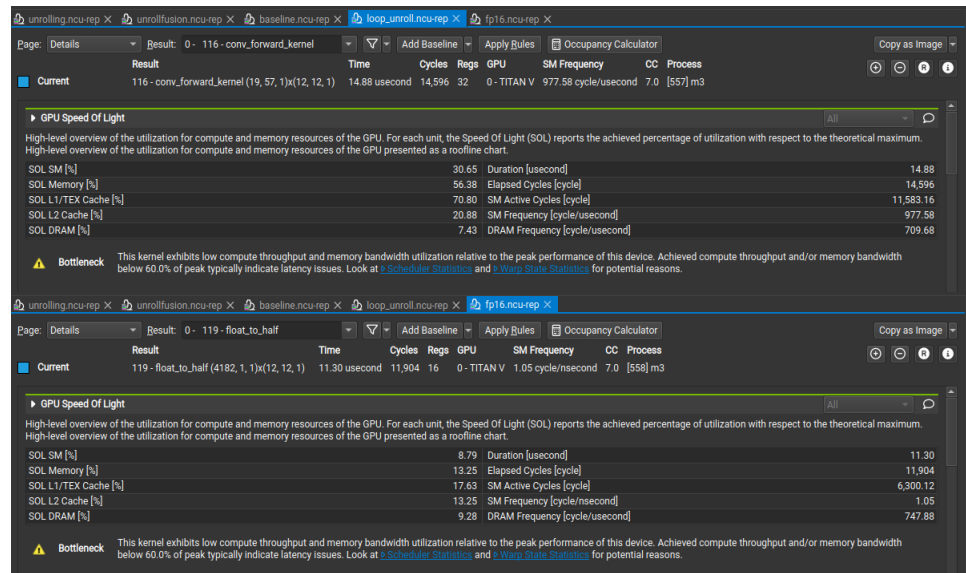
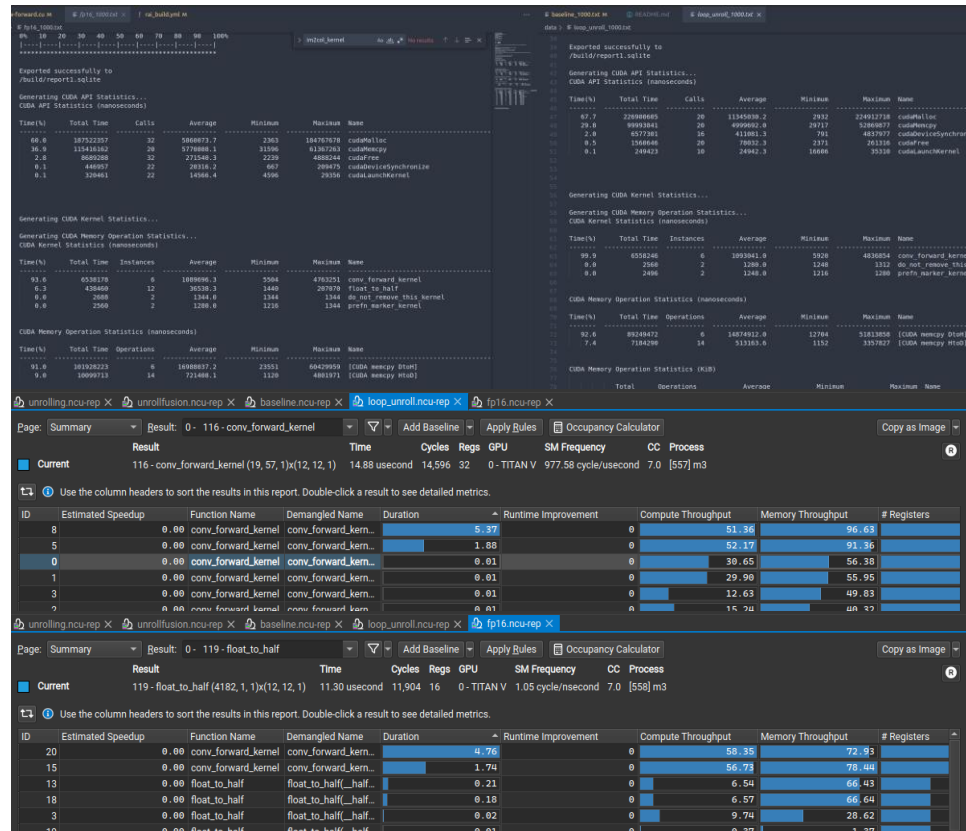
- b. How does the optimization work? Did you think the optimization would increase performance of the forward convolution? Why? Does the optimization synergize with any of your previous optimizations?

It works by transforming input matrix and mask matrix to FP16 accuracy. I think it will reduce the overhead of arithmetic calculation and reduce the memory throughput. This optimization is built from the loop unrolling optimization and is expected to reach the best performance.

- c. List the Op Times, whole program execution time, and accuracy for batch size of 100, 1k, and 5k images using this optimization (including any previous optimizations also used).

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.238143 ms	0.519883 ms	0m0.228s	0.86
1000	2.22948 ms	5.2218 ms	0m0.316s	0.887
5000	10.6534 ms	25.0527 ms	0m0.904s	0.8712

- d. Was implementing this optimization successful in improving performance? Why or why not? Include profiling results from *nsys* and *Nsight-Compute* to justify your answer, directly comparing to your baseline (or the previous optimization this one is built off of).



From result from NSYS and nsight compute we can see FP16 reduce the overhead of arithmetic operations as well as memory throughput as expected. This becomes my best implementation and is submitted with this version.

e. What references did you use when implementing this technique?

